

A Crash Course on Data Compression

2.1 Some Integer Codes in C++

Giulio Ermanno Pibiri

ISTI-CNR, giulio.ermanno.pibiri@isti.cnr.it



@giulio_pibiri



@jermmp

Overview

- C++ implementation of:
Unary, Gamma, Delta, Rice, Variable-Byte
- Compress/Uncompress some (long) integer lists
- Write-to/Load-from disk the compressed file

General Approach

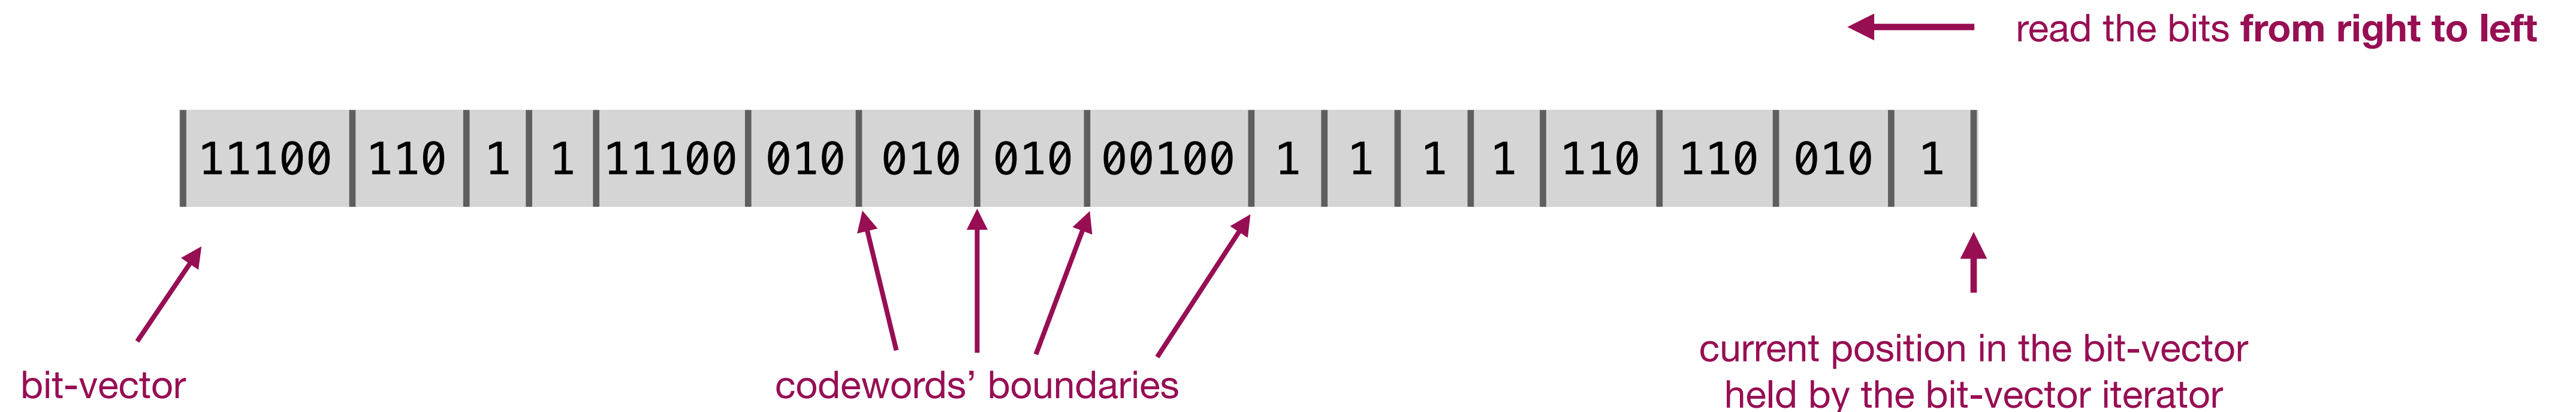
For the wanted code, implement **two** functions:

```
void write_codeword(bit_vector_builder& builder, uint64_t x)
```

which appends the codeword for the integer x to the bit-vector builder's bits, and

```
uint64_t read_codeword(bit_vector_iterator& it)
```

which decodes the codeword coming next the current position of the bit-vector's iterator.



Our Plan

```
/* Unary */
void write_unary(bit_vector_builder& builder, uint64_t x);
uint64_t read_unary(bit_vector_iterator& it);

/* Gamma */
void write_gamma(bit_vector_builder& builder, uint64_t x);
uint64_t read_gamma(bit_vector_iterator& it);

/* Delta */
void write_delta(bit_vector_builder& builder, uint64_t x);
uint64_t read_delta(bit_vector_iterator& it);

/* Rice */
void write_rice(bit_vector_builder& builder, uint64_t x, const uint64_t k);
uint64_t read_rice(bit_vector_iterator& it, const uint64_t k);

/* Variable-Byte */
void write_vbyte(bit_vector_builder& builder, uint64_t x);
uint64_t read_vbyte(bit_vector_iterator& it);
```

Our Plan

All you (essentially) need to implement our plan is:

the function

```
void bit_vector_builder::append_bits(uint64_t x, uint64_t len)
```

appends the `len` **least significant** bits
of `x` to the current end of the bit-vector



to implement `write_codeword`, and

the function

```
uint64_t bit_vector_iterator::take(uint64_t len)
```

reads the next `len` bits from the bit-vector
and returns them as the integer `x`



to implement `read_codeword`.

Unary


Represent $x \geq 0$ as $U(x) = 0^x 1$.

0 1
1 01
2 001
3 0001
...

```
void write_unary(bit_vector_builder& builder, uint64_t x) {  
    assert(x < 64);  
    uint64_t u = uint64_t(1) << x;  
    builder.append_bits(u, x + 1);  
}
```

```
uint64_t read_unary(bit_vector_iterator& it) {  
    return it.skip_zeros();  
}
```

remember: you read the bits **from right to left**, so you must first write/read the zeros



skip as many zeros as possible from the current position and return the number of skipped zeros



Delta

Write $b = |bin(x)| - 1$ using **Gamma**, followed by the b least significant bits of $bin(x)$.

```
void write_delta(bit_vector_builder& builder, uint64_t x) {  
    uint64_t xx = x + 1;  
    uint64_t b = msb(xx);  
    write_gamma(builder, b);  
    uint64_t mask = (uint64_t(1) << b) - 1;  
    builder.append_bits(xx & mask, b);  
}
```




← write in **Gamma** how many least significant bits we have

```
uint64_t read_delta(bit_vector_iterator& it) {  
    uint64_t b = read_gamma(it);  
    return (it.take(b) | (uint64_t(1) << b)) - 1;  
}
```

← read the **Gamma** code representing how many least significant bits we have

Rice

For a given parameter $k > 0$, write $q = \lfloor x/2^k \rfloor$ using **Gamma** followed by the reminder $r = x - q2^k$ in k bits.

```
void write_rice(bit_vector_builder& builder, uint64_t x, const uint64_t k) {  
    assert(k > 0);  
    uint64_t q = x >> k;   $\lfloor x/2^k \rfloor$   
    uint64_t r = x - (q << k);   $x - q2^k$   
    write_gamma(builder, q);  
    builder.append_bits(r, k);  
}  
  
uint64_t read_rice(bit_vector_iterator& it, const uint64_t k) {  
    assert(k > 0);  
    uint64_t q = read_gamma(it);  
    uint64_t r = it.take(k);  
    return r + (q << k);   $r + q2^k$   
}
```

Variable-Byte

Example for $x = 67822$, $\text{bin}(67822, 17) = 10000100011101110$.

00000100.10010001.11101110

recursive implementation
(for elegance)

```
void write_vbyte(bit_vector_builder& builder, uint64_t x) {  
    if (x < 128) {  
        builder.append_bits(x, 8);  
        return;  
    }  
    uint8_t data_bits = x & 127;  
    builder.append_bits(data_bits | 128, 8);  
    write_vbyte(builder, x >> 7);  
}
```

isolate the 7 data bits by masking
with 01111111 (127)

add the continuation bit by ORing
with 10000000 (128)

discard the 7 processed bits and
recurse on the rest

iterative implementation
(for efficiency)

```
uint64_t read_vbyte(bit_vector_iterator& it) {  
    uint64_t val = 0;  
    for (uint64_t shift = 0;; shift += 7) {  
        uint8_t byte = it.take_one_byte();  
        val += (byte & 127) << shift;  
        if (byte < 128) break;  
    }  
    return val;  
}
```

read one byte

get the 7 data bits and place
them into position by shifting

stop if control bit is 0

Compress

create the bit-vector's builder
reserve space for the number of lists

encode the gaps
between the integers

build the vector

write all the bytes
to the output file

```
template <typename WriteFunction>
void compress(std::string const& input_lists_filename,
             std::string const& output_filename,
             WriteFunction write)
{
    bit_vector_builder builder;
    builder.append_bits(0, 32);
    uint64_t num_lists = 0;

    std::ifstream in(input_lists_filename.c_str());
    while (!in.eof()) {
        uint64_t list_size = 0;
        in >> list_size;
        if (list_size > 0) {
            num_lists += 1;
            builder.append_bits(list_size, 32);

            uint32_t prev_x = 0;
            uint32_t x = 0;
            for (uint64_t i = 0; i != list_size; ++i) {
                in >> x;
                assert(x >= prev_x);
                write(builder, x - prev_x);
                prev_x = x;
            }
        }
        in.close();

        builder.set_bits(0, num_lists, 32);

        bit_vector bits;
        builder.build(bits);

        std::ofstream out(output_filename.c_str(), std::ofstream::binary);
        bits.save(out);
        out.close();
    }
}
```

input lists' filename
filename of the compressed output
accepts a write_codeword function

input is sorted

write the actual number
of compressed lists

example input: lists.txt

726
3
6
13
14
22
...
135
6
7
34
56
58
...

list size

726 lines

135 lines

```
void save(std::ofstream& out) const {
    out.write(reinterpret_cast<char const*>(&m_num_bits),
             sizeof(m_num_bits));
    out.write(reinterpret_cast<char const*>(m_bits.data()),
             m_bits.size() * sizeof(uint64_t));
}
```

Compress

```
int main(int argc, char** argv) {
    std::string type = argv[1];
    std::string input_lists_filename = argv[2];
    std::string output_filename = argv[3];

    if (type == "gamma") {
        compress(input_lists_filename, output_filename,
                [](bit_vector_builder& builder, uint32_t x) {
                    write_gamma(builder, x);
                });
    } else if (type == "delta") {
        compress(input_lists_filename, output_filename,
                [](bit_vector_builder& builder, uint32_t x) {
                    write_delta(builder, x);
                });
    } else if (type == "vbyte") {
        compress(input_lists_filename, output_filename,
                [](bit_vector_builder& builder, uint32_t x) {
                    write_vbyte(builder, x);
                });
    } else if (...) {
        ...
    } else {
        std::cout << "unknown type '" << type << "'" << std::endl;
        return 1;
    }

    return 0;
}
```

pass the specific write_codeword function



Decompress

```
void load(std::ifstream& in) {  
    in.read(reinterpret_cast<char*>(&m_num_bits), sizeof(m_num_bits));  
    m_bits.resize(num_64bit_words_for(m_num_bits));  
    in.read(reinterpret_cast<char*>(m_bits.data()),  
            m_bits.size() * sizeof(uint64_t));  
}
```

read all the bytes from
the file into memory

measure the decoding
time using `std::chrono`

```
template <typename ReadFunction>  
void decompress(std::string const& input_filename, ReadFunction read) {  
    typedef std::chrono::high_resolution_clock clock_t;  
    typedef std::chrono::microseconds duration_t;  
  
    bit_vector bits;  
  
    std::ifstream in(input_filename.c_str(), std::ifstream::binary);  
    bits.load(in);  
    in.close();  
  
    bit_vector_iterator it(bits);  
    uint64_t num_lists = it.take(32);  
    uint64_t num_ints = 0;  
  
    auto start = clock_t::now();  
    for (uint64_t i = 0; i != num_lists; ++i) {  
        uint64_t list_size = it.take(32);  
        uint32_t prev_x = 0;  
        uint32_t x = 0;  
        for (uint64_t i = 0; i != list_size; ++i) {  
            uint32_t x = read(it) + prev_x;  
            assert(x >= prev_x);  
            prev_x = x;  
        }  
        num_ints += list_size;  
    }  
    auto stop = clock_t::now();  
    auto elapsed = std::chrono::duration_cast<duration_t>(stop - start);  
  
    std::cout << "decompressed " << num_ints << " integers in "  
              << elapsed.count() << " microsecs" << std::endl;  
    std::cout << "(" << (elapsed.count() * 1000.0) / num_ints << " ns/int)"  
              << std::endl;  
}
```

accepts a `read_codeword` function

read the number of lists

decode each list

decode each gap

Decompress

```
int main(int argc, char** argv) {
    std::string type = argv[1];
    std::string input_filename = argv[2];

    if (type == "gamma") {
        decompress(input_filename,
            [](bit_vector_iterator& it) { return read_gamma(it); });
    } else if (type == "delta") {
        decompress(input_filename,
            [](bit_vector_iterator& it) { return read_delta(it); });
    } else if (type == "vbyte") {
        decompress(input_filename,
            [](bit_vector_iterator& it) { return read_vbyte(it); });
    } else if (...) {
        (...)
    } else {
        std::cout << "unknown type '" << type << "'" << std::endl;
        return 1;
    }

    return 0;
}
```

pass the specific read_codeword function



How to compile and run the code

From a terminal window, move into this folder and type the following commands.

To compile in a "debug" environment, define (`-D`) the `DEBUG` flag to enable all asserts:

```
g++ -std=c++11 -DDEBUG compress.cpp -o compress
g++ -std=c++11 -DDEBUG decompress.cpp -o decompress
g++ -std=c++11 -DDEBUG check.cpp -o check
```

To compile for maximum speed, disable all asserts (`-DNDEBUG`) and also use the optimization flags `-O3` and `-march=native` :

```
g++ -std=c++11 -DNDEBUG -O3 -march=native compress.cpp -o compress
g++ -std=c++11 -DNDEBUG -O3 -march=native decompress.cpp -o decompress
g++ -std=c++11 -DNDEBUG -O3 -march=native check.cpp -o check
```

Now, first unzip the file `lists.txt.gz` which contains 10 sorted integer lists:

```
gunzip lists.txt.gz
```

Then use the program `./compress` to actually compress the lists. The program expects the following arguments:

```
Usage: ./compress [type] [input_lists_filename] [output_filename]
```

where `type` is one of the following: `gamma` , `delta` , `vbyte` , `rice_k1` , or `rice_k2` .

Below, some examples with type `gamma` .

```
./compress gamma lists.txt out_gamma.bin
./decompress gamma out_gamma.bin
./check gamma out_gamma.bin lists.txt
```

The script `run_all.sh` shows all the examples. To run it, use:

```
bash run_all.sh
```

Micro benchmark

On a desktop Mac book pro (16-inch, 2019) with a 2.6 GHz 6-Core Intel Core i7 processor, I got the following results (compiling with optimization flags `-O3` and `-march=native`).

Code	bits/int	ns/int
gamma	4.24	3.89
delta	4.96	4.84
rice k=1	3.49	4.77
rice k=2	3.89	4.77
vbyte	8.11	0.95